


# **MUSIC GENRE CLASSIFICATION**

Using Machine Learning Techniques




# TABLE OF CONTENTS

1. Problem Statement
2. Data-Set
3. Data Pre-processing
4. Models Explored
5. Results and Discussion
6. Conclusion and Future Work

Honour Code:

We hereby affirm that we have upheld the highest principles of honesty and integrity while we worked on this project.



# **GROUP MEMBERS**

MANAS JAIN (170040068)

NIKHIL CHOUDHARY (170110048)

HARSHITA MASAND (170110064)

# PROBLEM STATEMENT


Music Genre Classification  
using Machine Learning and  
Deep Learning Techniques

Genre classification is an important task with many real world applications. As the quantity of music being released on a daily basis continues to sky-rocket the need for accurate meta-data required for database management and search/storage purposes climbs in proportion.

# INTRODUCTION

Categorizing music files according to their genre is a challenging task in the area of music information retrieval (MIR). In this project, we implemented four different classification algorithms admitting the input - transformed mel-spectrograms of our raw amplitude data. We implemented the following algorithms :


- K-Nearest Neighbours
- Support Vector Machine
- Feed Forward Neural Network
- Convolutional Neural Network



# GENERAL ALGORITHM




Audio File



Mel Spectrogram


Feature  
Extraction and  
classification

# BACKGROUND




**2002**

The mixture of Gaussians model and k-nearest neighbors was used. It achieved 61% accuracy [1]



**2003**

Multiple layers of SVM achieved upto 90% accuracy. [2]



**2012 - PRESENT**

CNN prove to be incredibly accurate music genre classifiers with an accuracy of around 91% [3]

# DATASET

We used labeled public GTZAN dataset (1k 30s audio clips). From each clip, we sampled 2s window at 8 random location, thus augmenting our data to 8000 clips of two seconds each. Data was sampled at 22050Hz, this leaves us with 44100 features.

From Kaggle, we imported 2 CSV files - One file has for each song (30 seconds long) a mean and variance computed over multiple features that can be extracted from an audio file. The other file has the same structure, but the songs were split before into 3 seconds audio files (this way increasing 10 times the amount of data we fuel into our classification models).

Training Data	8000
Testing Data	100
Cross-Validation	100

No. of samples*No. Of features	9990*58
Testing data	3296 (33%)
Training data	6694 (67%)




# PRE - PROCESSING

- Converting the raw audio into mel-spectrograms
- Mel-spectrograms are a commonly used method of featurizing audio because they closely represent how humans perceive audio i.e. in log frequency
- For this one must apply short-time Fourier transforms across sliding windows of audio, most commonly around 20ms wide

$$\text{STFT}\{x[n]\}(m, \omega) = \sum x[n]w[n - m]e^{-j\omega n}.$$

where signal  $x[n]$ , window  $w[n]$ , frequency axis  $\omega$ , and shift  $m$  computed using sliding DFT algorithms




# PRE - PROCESSING


- STFT are then mapped to the mel scale by transforming the frequencies  $f$  by

$$m = 2595 \log_{10}\left(1 + \frac{f}{700}\right).$$


- Then we take the discrete cosine transform of the result (common in signal processing) in order to get our final output mel-spectrogram
  - We used Librosa library and chose to use 64 mel-bins and a window length of 512 samples with an overlap of 50% between windows
- Shape of our data (8000,64,173)**




# SAMPLE MEL SPECTROGRAMS OF OUR PRE-PROCESSED DATA




Blues




Classical




Country



Rock



Jazz



Pop

# MODELS IMPLEMENTED

## SVM

We created a model with RBF Kernel. This gave a baseline accuracy to compare with the Deep-Learning Models.



## K-NN

Post PCA, we used  $k=10$  and distance weighting.

## FEED FORWARD NEURAL NETWORK

Our network has 3-fully connected layers, each using ReLU activation. We use softmax o/p with cross-entropy loss and Adam optimisation.



## CNN


Adam optimisation with ReLU Activation.

# PRINCIPAL COMPONENT ANALYSIS

- PCA to reduce dimension for two of our models (k-NN & SVM)
- In order to do this, we first transformed the mel-spectrograms by normalizing their mean and variance

$$\frac{1}{m} \sum_{i=1}^m (x^{(i)T} u)^2 = \frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T}.$$

- We had best results reducing to 15 dimensions, analogous to taking the top 15 eigenvectors of  $\Sigma$  and projecting our data onto the subspace spanned by them. We implemented this using scikit-learn in python.




# K-NN

- Since there are 10 genres we got best accuracy from K=10
- Those which return the largest value on  $\|x - x^{(i)}\|_2$ , where  $\|\cdot\|^2$  denotes the Euclidean distance between points, be the 10 closest neighbours to x.
- We choose weights  $w_i$  for each i such that:

$$w_i \propto \frac{1}{\|x - x^{(i)}\|_2}, \quad \sum_{i=1}^K w_i = 1.$$


$$\arg \max_y \sum w_i \mathbb{1}(y = y^{(i)}),$$



- Finally we return label which is most prevalent among x's 10 nearest neighbors when weighted by distance. This was implemented using scikit-learn in python.

# K-NN - RESULTS

Accuracy on Training Data	1
Accuracy on CV	0.53
Accuracy on Testing Data	0.54



# SUPPORT VECTOR MACHINE


To make our data linearly separable, we used the following kernel function:

$$\min_{\gamma, w, b} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^m \xi_i \quad \text{s.t. } y^{(i)} \left( \sum_j \alpha_j K(x^{(j)}, x^{(i)}) + b \right) \geq 1 - \xi_i \text{ and } \xi_i \geq 0$$

where  $x^{(i)}$  are examples,  $\alpha_j$ ,  $b$  are weights and biases,  $C$  is a penalty parameter, and  $1 - \xi_i$  is the functional margin for example  $i$ .  $K: \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$  is a kernel function. We used the RBF Kernel Function:

$$K(x^{(j)}, x^{(i)}) = \exp \left( -\frac{\|x^{(j)} - x^{(i)}\|_2^2}{2\sigma^2} \right)$$


This kernel, also sometimes called the Gaussian kernel, corresponds to an infinite dimensional feature space is related to Euclidean distance. This function as a whole is often minimized using sequential minimization optimization.





# SVM - RESULTS

Accuracy on Training Data	0.59025
Accuracy on CV	0.5
Accuracy on Testing Data	0.5



This model was implemented using Scikit-Learn in Python.

# FEED - FORWARD NEURAL NETWORK

We used a fully connected neural network as well, with ReLU activation and 6 layers, with cross-entropy loss. As the input to our model was 1D, when using mel-spectrograms, we flattened the data. Our model is fully connected, which means each node is connected to every other node in the next layer. At each layer, we applied a ReLU activation function to the output of each node, following the formula:

$$\text{ReLU}(x) = \begin{cases} x & x \geq 0 \\ 0 & x < 0. \end{cases}$$


we construct a probability distribution of the 10 genres by

$$\sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

To optimize our model, we minimized cross entropy loss:

$$CE(\theta) = - \sum_{x \in X} y(x) \log \hat{y}(x)$$

This was implemented with TensorFlow



# FEED FORWARD NN - ARCHITECTURE


The model has been implemented using Keras and Tensor Flow.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 256)	15104
=====		
dense_1 (Dense)	(None, 128)	32896
=====		
dense_2 (Dense)	(None, 64)	8256
=====		
dense_3 (Dense)	(None, 10)	650
=====		
Total params: 56,906		
Trainable params: 56,906		
Non-trainable params: 0		
=====		



# FEED FORWARD NN - RESULTS

Accuracy on Training Data	0.9994
Accuracy on Testing Data	0.8893



# CONVOLUTIONAL NEURAL NETWORK

This was our most advanced model, using 3 convolution layers, each with its own max pool and regularization, feeding into 3 fully connected layers with ReLU activation, softmax output, and cross entropy loss.



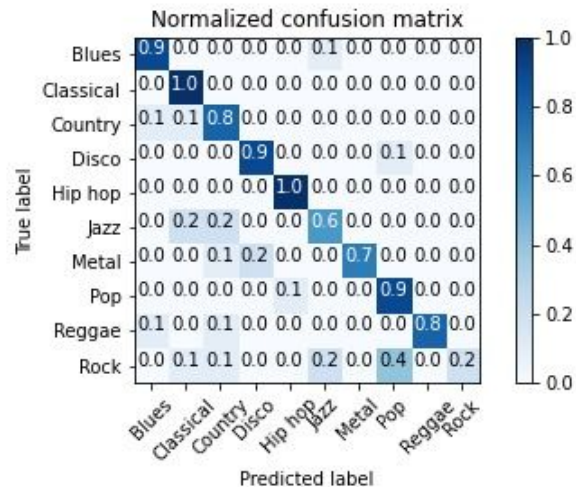
# CONVOLUTIONAL NEURAL NETWORK- IMPLEMENTATION DETAILS

As the accuracy in our Cross-Validation crossed 80%, the model saved the confusion matrix. This happened for us at No. of Epochs=150, Batch size:=200

Model: "sequential_2"			dropout_7 (Dropout)	(None, 6, 9, 64)	0
Layer (type)	Output Shape	Param #	flatten_2 (Flatten)	(None, 3456)	0
=====			dense_6 (Dense)	(None, 64)	221248
conv2d_6 (Conv2D)	(None, 61, 170, 64)	1088	dropout_8 (Dropout)	(None, 64)	0
batch_normalization_4 (Batch Normalization)	(None, 61, 170, 64)	256	dense_7 (Dense)	(None, 32)	2080
max_pooling2d_6 (MaxPooling2D)	(None, 30, 42, 64)	0	dense_8 (Dense)	(None, 10)	330
conv2d_7 (Conv2D)	(None, 28, 38, 64)	61504	=====		
max_pooling2d_7 (MaxPooling2D)	(None, 14, 19, 64)	0	Total params: 303,210		
dropout_6 (Dropout)	(None, 14, 19, 64)	0	Trainable params: 302,954		
conv2d_8 (Conv2D)	(None, 13, 18, 64)	16448	Non-trainable params: 256		
batch_normalization_5 (Batch Normalization)	(None, 13, 18, 64)	256			
max_pooling2d_8 (MaxPooling2D)	(None, 6, 9, 64)	0			

# CONVOLUTIONAL NEURAL NETWORK- RESULTS


Training Accuracy	0.932250
Cross-Validation Accuracy	0.843750
Testing Accuracy	0.828125



This model was implemented using tensorflow and keras in Python. Google Colab free GPU was used for training the model.

# RESULTS AND DISCUSSION


- Since many rock music excerpts lack the easily visible beats that other genres such as hip-hop and disco possess, we found by looking at our confusion matrix that our CNN managed to correctly classify 50% of rock audio as rock, labeling the others as mainly country or blues. Additionally, it incorrectly classified some country, as well as a small fraction of blues and reggae, as rock music.
- We also noted that correct classification of jazz was less accurate than most other categories. Our algorithm falsely classified some jazz music as classical, although never did the reverse of this.
- The accuracy achieved by FFNN is 88%, which is also good given the huge amount of data we have. However, it requires feature extraction and used a large number of features that made this computationally heavy. Hence, CNN is preferred.






# CONCLUSION & FUTURE WORK


- CNN performed the best, as we expected. It took the longest time to train as well, but the increase in accuracy justifies the extra computation cost.
- In the future, given that this is time series data, some sort of RNN model may work well (GRU, LSTM)
- Also we may have opportunities for transfer learning, for example in classifying music by artist or by decade.




# CONTRIBUTIONS



**MANAS JAIN**  
CNN+ SVM



**HARSHITA MASAND**  
k-NN+ PPT



**NIKHIL  
CHOUDHARY**  
Feed Forward Neural  
Network+ PPT

# THANK YOU !

**Please keep this slide for attribution**

CREDITS: This presentation template was  
created by **Slidesgo**, including icons by  
**Flaticon**, infographics & images by **Freepik**



# REFERENCES

1. G. Tzanetakis and P. Cook. Musical genre classification of audio signals. IEEE Transactions on Speech and Audio Processing, 10(5):293–302, July 2002.
2. G. Tzanetakis and P. Cook. Musical genre classification of audio signals. IEEE Transactions on Speech and Audio Processing, 10(5):293–302, July 2002.
3. Y. Panagakis, C. Kotropoulos, and G. R. Arce. Music genre classification via sparse representations of auditory temporal modulations. In 2009 17th European Signal Processing Conference, pages 1–5, Aug 2009.